

Introduction to R

Anindya Bhattacharya, PhD | abhattacha3@ucla.edu
CBI fellow | Xia Yang lab

Anindya Bhattacharya, PhD

abhattacha3@ucla.edu

Office hours: 2000 TLSB

Xia Yang lab

Department of Integrative Biology and Physiology



Areas of interest: miRNA, genetic variation, gene expression data analyses, machine learning,

Collaboratory Website

<http://collaboratory.lifesci.ucla.edu/>



Workshop 3: Introduction to R

▶ Day 1

- ▶ Introduction to R
 - ▶ Installation, configuration, basic concepts, syntax and usage

▶ Day 2

- ▶ Data Analysis
 - ▶ Visualization methods
 - ▶ Statistical methods

▶ Day 3

- ▶ Scripting, batch mode
- ▶ Packages, Bioconductor



Workshop 3: Introduction to R

▶ **Day I**

- ▶ What is R? Why use it?
- ▶ Getting up and running
 - ▶ Installation, configuration issues
 - ▶ overview of R interface
 - ▶ GUIs for working with R
- ▶ Your first R session
- ▶ Basic concepts
 - ▶ Starting and stopping an R session
 - ▶ finding help/documentation
 - ▶ Data structures
 - ▶ Getting data in and out of R
 - ▶ Basic analysis



Dropbox

- ▶ In the drop box folder are:
 - ▶ course powerpoint
 - ▶ sample scripts
 - ▶ exercises
- ▶ Please **DO NOT** make changes to these files because they will affect **EVERYONE!**
- ▶ **INSTEAD**, copy them to a directory on your computer
- ▶ Let me know if you need to be invited to the dropbox folder



What is R?

- ▶ A glorified calculator
- ▶ Statistical tables
- ▶ Visualization tool
- ▶ Programming language and environment



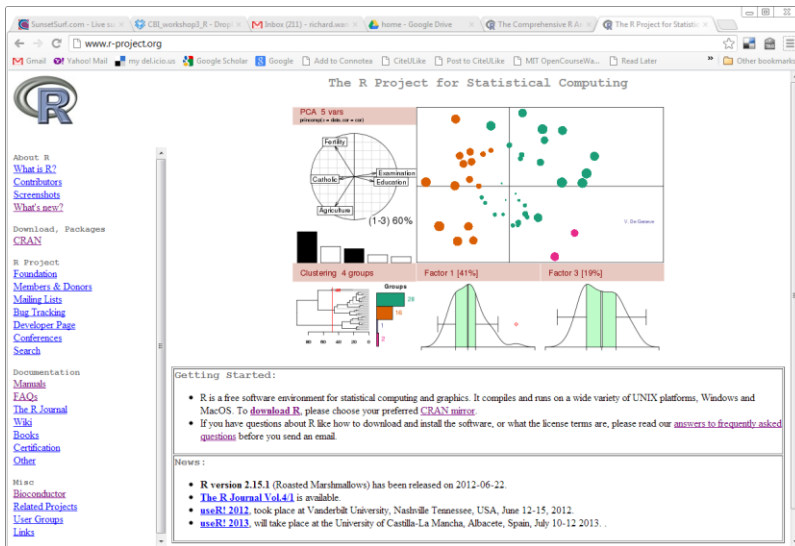
Why R?

- ▶ Many tools STATA, SPSS, Matlab, Excel...
- ▶ Advantages
 - ▶ powerful
 - ▶ up to date with latest algorithms (packages, Bioconductor)
 - ▶ strong community of users
 - ▶ Preferred by statistics community
 - ▶ it's free
- ▶ Disadvantages
 - ▶ steep learning curve, but can be useful quickly
 - ▶ not pretty
 - ▶ can be memory intensive



Where to download R?

- Use <http://r-project.org> or a CRAN mirror (<http://cran.stat.ucla.edu>)



The R Project for Statistical Computing

PCA: 5 vars (Principal Component Analysis)

Clustering: 4 groups

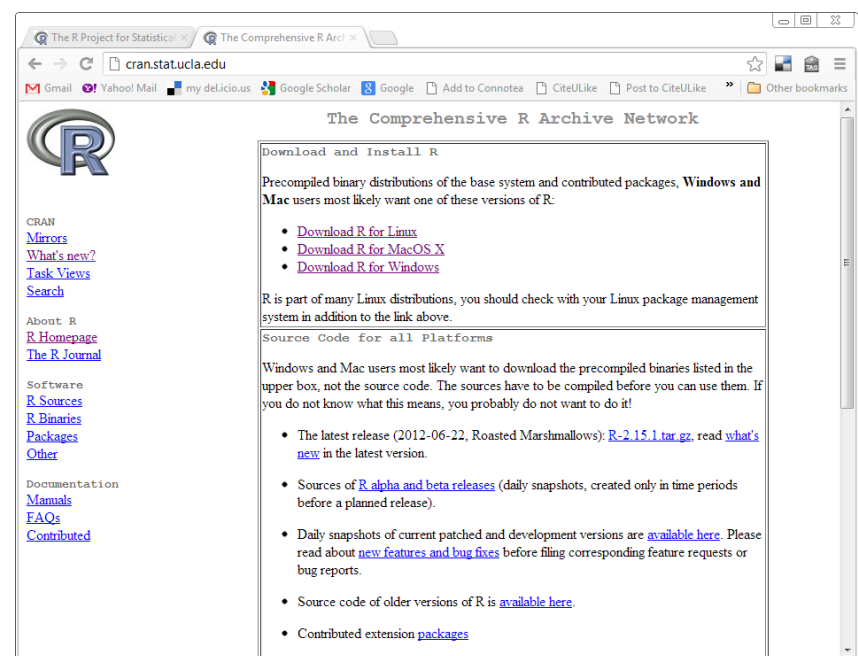
Factor 1 (41%) Factor 3 (19%)

Getting Started:

- R is a free software environment for statistical computing and graphics. It compiles and runs on a wide variety of UNIX platforms, Windows and MacOS. To [download R](#), please choose your preferred [CRAN mirror](#).
- If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.

News:

- R version 2.15.1 (Roasted Marshmallows) has been released on 2012-06-22.
- The R Journal Vol.4.1 is available.
- ruseR! 2012, took place at Vanderbilt University, Nashville Tennessee, USA, June 12-15, 2012.
- ruseR! 2013, will take place at the University of Castilla-La Mancha, Albacete, Spain, July 10-12 2013.



The Comprehensive R Archive Network

Download and Install R

Precompiled binary distributions of the base system and contributed packages, **Windows** and **Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for MacOS X](#)
- [Download R for Windows](#)

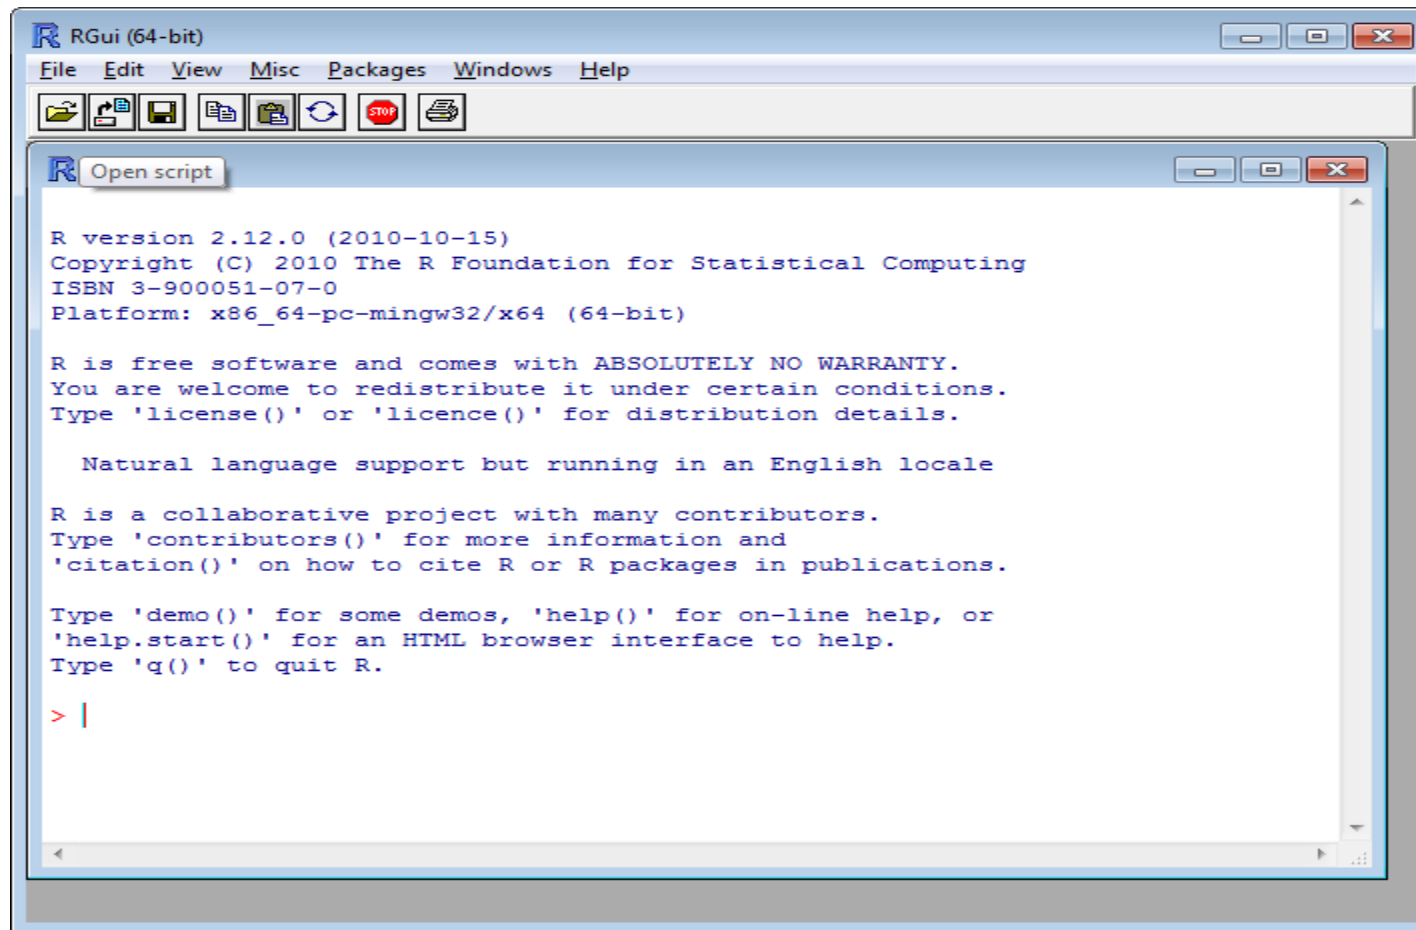
R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.

Source Code for all Platforms

Windows and Mac users most likely want to download the precompiled binaries listed in the upper box, not the source code. The sources have to be compiled before you can use them. If you do not know what this means, you probably do not want to do it!


- The latest release (2012-06-22, Roasted Marshmallows): [R-2.15.1.tar.gz](#), read [what's new](#) in the latest version.
- Sources of [R alpha and beta releases](#) (daily snapshots, created only in time periods before a planned release).
- Daily snapshots of current patched and development versions are [available here](#). Please read about [new features](#) and [bug fixes](#) before filing corresponding feature requests or bug reports.
- Source code of older versions of R is [available here](#).
- Contributed extension [packages](#)

Starting Up...



Stopping

- ▶ To end your session
 - ▶ Close window or `q()`
- ▶ What does this mean?

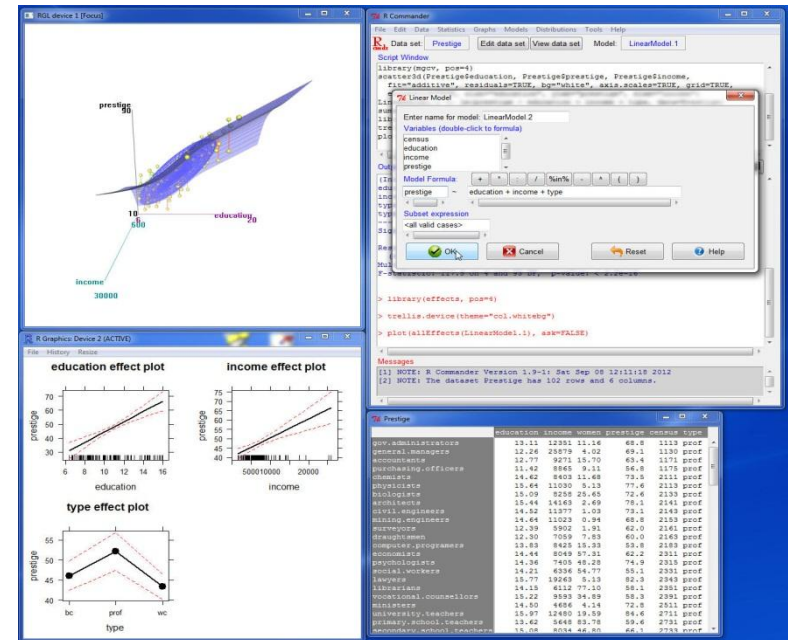


- ▶ It saves everything you created (variables, etc) into a (hidden) file in the current directory usually named `.Rdata`

GUIs and IDEs

- ▶ R interface is pretty sparse
- ▶ Some alternatives:

	Win	Mac	Linux
RStudio Desktop	x	x	x
RStudio Server			x
Rcommander	x	x	x
Tinn-R	x		
JGR	x	x	x
RKward	x		x
Emacs/ESS	x	x	x
Vi	x	x	x
notepad	x	x	x



RStudio

File Edit Code View Plots Session Build Debug Tools Help

Go to file/function

mutation.R x mutation_WCage.R x mutation.R x R_refresher.Rmd x Untitled1* x

Source on Save Run Source

```
1 ### our code goes here
2
3 x = read.table("genes.fpk_tracking", header=T)
```

3:47 (Top Level) R Script

Console

```
> rats
  id  sex weight length
1 rat1 female    2   100
2 rat2 female    4   105
3 rat3 female    1   115
4 rat4 female   11   130
5 rat5 female   18    95
6 rat6 male    12   150
```

Studio Home RStudio IDE Shiny Training Projects About Blog

Welcome to RStudio

Open source and enterprise-ready professional software for the R community



Powerful IDE for R
RStudio IDE is a powerful and productive user interface for R. It's free and open source, and works great on Windows, Mac, and Linux.
[Download now](#) [Learn more](#)

Web framework for R
Shiny is an elegant and powerful web framework for building interactive reports and visualizations using R — with or without web development skills.
[Get started](#)

Open source R packages
Our developers and expert trainers are the authors of several popular R packages, including ggplot2, plyr, lubridate, and others.
[See projects](#)

It's like unix

- ▶ Many of the commands are similar to command line unix tools (cat, ls, grep, ...)
- ▶ but in R, everything is a function so **()** is important!
 - ▶ cat()
 - ▶ ls()
 - ▶ rm()
- ▶ What happens if you forget the parentheses?
 - ▶ try it!



R packages

► List installed packages

```
> installed.packages()  
> a<- installed.packages()
```

► Available packages

```
> p<-available.packages()  
> dim(p)
```

► Is the package “boot” already installed?

```
> is.element("boot", a[,1])  
>[1] TRUE
```



Finding help

```
> help("data.frame")
> ?data.frame
> help.search("chisquare")
> apropos("mean")
[1] "colMeans"          "kmeans"            "mean"              "mean.data.frame"
[5] "mean.Date"         "mean.default"      "mean.difftime"     "mean.POSIXct"
[9] "mean.POSIXlt"      "rowMeans"          "weighted.mean"
```



```
> library(help=utils)
> help(BATCH)
```



Advice for learning R

- ▶ There are no easy ways to learn R except by
 1. Typing examples in yourself! It makes a difference (muscle memory, syntax)
 2. Persisting!
 3. Iterating!
- ▶ Steep learning curve, but with a few commands it can be useful quickly
- ▶ Try things out – you will not break your data!



Getting around

- ▶ **R starts in a particular directory, usually the one where R is located**

```
# print the current working directory
getwd()

# set the current working directory
setwd("/home/.../project1")

# show the files in your current directory
dir()
```

- ▶ **Create a directory for your project**

- ▶ **good way to organize**

```
>dir.create("R_CODE")
```



A few things

- ▶ comments start with a hash tag (#) and are ignored by R. Great way to make notes!
- ▶ The prompt ">" is where you type commands, statements

```
> a = 1
```

- ▶ "+" is a continuation character; if your statement is not complete, R prompts you to finish it:

```
> a =
```

```
+ 3
```

```
> a
```

```
[1] 3
```

- ▶ Typing the variable name gives you its contents

```
> a
```

```
[1] 3
```

- ▶ Case matters! **A** is not the same as **a**

```
> A
```

```
Error: object 'A' not found
```

- ▶ Typing a function name without parens () gives the source code



Try it out

```
# genes.fpkms_tracking output of cufflinks
# it is tab delimited with a header line
# what happens if we say header=F?
x = read.table("genes.fpkms_tracking", header=T)
ls()
# what do the numbers in dim() mean?
dim(x)
names(x)
head(x)
class(x)
sum(x[,11] > 100)
sum(x$q1_FPKM > 100)
# try changing the axis limits "ylim" and printing character "pch"
plot(x[,11], x[,14], xlim=c(0, 10), ylim=c(0,20), xlab="exp", main="my
experiment", pch=2)
# try changing the number of breaks
hist(x[,11], col="red", breaks=100)
cor(x[,11], x[,14])
newx = x[ x$status=="OK", ]
ls()
```



Basic data structures in R

	Homogeneous	Heterogeneous
1d	Vector	List
2d	Matrix	Data frame
nd	Array	

Note: R has no 0-dimensional, or scalar types. Individual numbers or strings, which you might think would be scalars, are actually vectors of length one.



Indexing

▶ let's create a vector

```
a = c(1,2,3,4)
a[2] # gives 2, indexing in R start with 1 not 0!
a[1:3]
```

▶ create a matrix

```
m = matrix(c(1,2,3,4,5,6), nrow=2, ncol=3)
m[,1]
[1] 1 2
m[,2] = 30
```

▶ logical indexing

```
x = 1:10
x > 5
[1] FALSE FALSE FALSE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE
x[ x > 5]
[1]  6  7  8  9 10
```



Vectors

► Basically arrays

```
> v = 100
```

```
> v <- 100
```

```
> c(1, 2, 3)
```

```
> 1:3
```

```
> rep(1:2, 3)
```

```
> vector(mode="numeric", length=5)
```

► Vector Arithmetic

```
> 1 + 1:3
```

```
[1] 2 3 4
```

```
> 1:2 + 1:3 #what happens?
```



Vector Arithmetic

► Vector Arithmetic

```
> 1:2 + 1:3 #what happens? - ERROR!
```

```
> 1:2 + 1:10
```

```
[1]  2  4  4  6  6  8  8 10 10 12
```

► Vector indexing (1-based)

```
> b = 1:2 + 1:10
```

```
> b[1]
```

► Vector length

```
> length(b)
```



Name

► **You can name a vector in three ways:**

```
> x <- c(a = 1, b = 2, c = 3)
```

```
> x <- 1:3; names(x) <- c("a", "b", "c")
```

```
> x <- setNames(1:3, c("a", "b", "c"))
```



Factors : Categorical data

► Categorical variables

► not the same as character class

```
> col = c("red", "green", "red", "yellow", "red")
> class(col)
[1] "character"
> col.f = factor(col)
[1] red      green    red      yellow   red
Levels: green red yellow
> levels(col.f) = c("green", "red", "blue")
> col.f
[1] red    green  red    blue   red
Levels: green red blue
> table(col.f)
col.f
green  red  blue
     1   3   1
```



Matrices

► Grid of numbers. Must be numeric!

```
> x <- matrix(1:10, nrow=2)
> dim(x)
[1] 2 5
> x
[,1] [,2] [,3] [,4] [,5]
[1,]      1      3      5      7      9
[2,]      2      4      6      8     10
> as.vector(x)
[1] 1 2 3 4 5 6 7 8 9 10
> dim(x)
[1] 2 5
```

hint: `x[2,4]` means get element from row 2, col 4



List

- ▶ Lists are different from atomic vectors because their elements can be of any type, including lists. You construct lists by using `list()` instead of `c()`.

```
# define list x
>x <- list(1:3, "a", c(TRUE, FALSE, TRUE), c(2.3, 5.9))
# display structure of x
>str(x)
# vector from a list
>unlist(x, recursive = TRUE, use.names = TRUE)
```



Matrices and arrays

- ▶ Adding a `dim()` attribute to an atomic vector allows it to behave like a multi-dimensional array.

Two scalar arguments to specify rows and columns

```
>a <- matrix(1:6, ncol = 3, nrow = 2)
```

```
>rownames(a) <- c("A", "B")
```

```
>colnames(a) <- c("a", "b", "c")
```

One vector argument to describe all dimensions

```
>b <- array(1:12, c(2, 3, 2))
```

```
>dimnames(b) <- list(c("one", "two"), c("a", "b", "c"),  
  c("A", "B"))
```



Data frames

- ▶ the most useful data structure
- ▶ data frames hold data like a spreadsheet
 - ▶ observations are rows
 - ▶ covariates are columns
- ▶ like matrices, they can be subsetted and indexed with two subscripts



Data frames

```
> df = data.frame(type=rep(c("case", "control"), c(2, 3)),  
time=rexp(5))
```

```
> df
```

	type	time
1	case	1.1745712
2	case	1.1691266
3	control	0.8227643
4	control	0.1301390
5	control	1.0581316

```
> df$time
```

```
[1] 1.1745712 1.1691266 0.8227643  
[4] 0.1301390 1.0581316
```

```
> df["time"]
```



Data frames

```
# let's update row names
> names(df)
[1] "type" "time"
> rn <- paste("id", 1:5, sep="")
> rownames(df) <- rn
> df[1:2, ]
      type  time
id1 case  1.174571
id2 case  1.169127
# update column names
> names(df) = c("var1", "var2")
> df[1:2,]
      var1 var2
id1 case 0.1653743
id2 case 0.4509330
```



Data type conversions

- ▶ Many problems are due to unexpected data types
 - ▶ eg. using a factor in place of a character string
- ▶ to convert from one data type to another, try

```
as.vector()
```

```
as.character()
```

```
as.data.frame()
```

```
as.numeric()
```

```
as.matrix()
```

- ▶ to determine what data type your variable is, try

```
class(variable)
```



Functions

▶ you can define your own functions

```
> square = function(x) {  
+           x*x  
+           }  
  
> square(10)  
[1] 100
```

```
> square(1:4)  
[1] 1 4 9 16
```

▶ R functions take can take named and default parameters

```
say = function(word="hi", repeats=3) {  
  cat( rep(word, repeats), "\n")  
}  
  
say()  
say(word="hola", repeats=1)  
say(repeats=2)
```



Reading in data

- ▶ A very common operation

```
> x = read.table("myfile.txt")
```

- ▶ Several variants: `read.csv()`, `read.delim()`
- ▶ this reads a delimited file into a `data.frame` object
- ▶ every line of the file must have the same number of elements!
 - ▶ sometimes you can fix things eg `read.table("myfile.txt", sep="\t")`
- ▶ entire file read into memory, so this can be slow but optimizations possible



Reading the help

```
read.table(file, header = FALSE, sep = ",", quote = "\"\"",  
           dec = ".", row.names, col.names,  
           as.is = !stringsAsFactors,  
           na.strings = "NA", colClasses = NA, nrows = -1,  
           skip = 0, check.names = TRUE, fill = !blank.lines.skip,  
           strip.white = FALSE, blank.lines.skip = TRUE,  
           comment.char = "#",  
           allowEscapes = FALSE, flush = FALSE,  
           stringsAsFactors = default.stringsAsFactors(),  
           fileEncoding = "", encoding = "unknown")
```

```
read.csv(file, header = TRUE, sep = ",", quote="\"", dec=".",  
         fill = TRUE, comment.char="", ...)
```

```
read.csv2(file, header = TRUE, sep = ";", quote="\"", dec=".",  
          fill = TRUE, comment.char="", ...)
```

```
read.delim(file, header = TRUE, sep = "\t", quote="\"", dec=".",  
           fill = TRUE, comment.char="", ...)
```

```
read.delim2(file, header = TRUE, sep = "\t", quote="\"", dec=".",  
            fill = TRUE, comment.char="", ...)
```



Writing data

- ▶ Writing data to a file is pretty simple
- ▶ Some odd defaults need to be commonly overridden

```
write.table(x, file="file_name.txt", quote=F, sep="\t",  
row.names=F, col.names=T)
```

- ▶ Other variants `write.csv()`



Saving and loading session

- ▶ all objects in your workspace can be saved to a file that can be read again on the next startup
- ▶ `save()` or `save.image()`
- ▶ the file is usually called `.RData` in your current directory
- ▶ `load()` will load an `.RData` file
- ▶ GUI might have these options



Some useful commands

- ▶ show a recent history of commands you typed
 - ▶ `history()`
- ▶ customizing your R
 - ▶ `.Rprofile` is a file found in your home directory where you can set defaults, define some commonly used functions, etc
- ▶ `rm()` – remove an object from workspace
- ▶ `ls()` – list the objects in your workspace
- ▶ `cbind(array1, array2)` – combine columns
- ▶ `rbind(array1, array2)` – combine rows
- ▶ `t()` – transpose a vector, matrix or data frame
- ▶ `sessionInfo()` – info about your settings



Match

- ▶ **match** returns a vector of the positions of (first) matches of its first argument in its second.

```
> df1 = read.table("df1.txt", header=T)
> df2 = read.table("df2.txt", header=T)
> match(df2$id, df1$id)
[1] 1  4  8  9 NA NA
```

```
# alternate syntax
> df2[,1] %in% df1[,1]
[1] TRUE TRUE TRUE TRUE FALSE FALSE
```



Merge

► Merge two data frames that have a common identifier

```
df1 = read.table("df1.txt", header=T)
df2 = read.table("df2.txt", header=T)
both = merge(df1, df2, by.x=1, by.y=1)
new = merge(df1, df2, by.x=1, by.y=1, all.x=T)
```



which

- ▶ **which** helps you find the index of an element

```
> a = 1:10
```

```
> which(a < 5)
```

```
[1] 1 2 3 4
```

```
> a[which(a < 5)]
```

```
[1] 1 2 3 4
```

```
> ids = which(a < 5)
```

```
[1] 1 2 3 4
```

```
> a[ids]
```



Special values

- ▶ **NA (Not available)**
 - ▶ sometimes data shows up as 'NA'
 - ▶ usually due to
 - ▶ missing data
 - ▶ conversion
 - ▶ **not** the same as zero (0) !
 - ▶ test with `is.na(x)`
- ▶ **NaN (not a number)**
 - ▶ `0/0`
 - ▶ `is.nan()`
- ▶ **Inf (infinite) or -Inf**
 - ▶ `10/0`
 - ▶ test with `is.finite(x)` or `is.infinite(x)`



NA/NAN/INF

```
#Test if any NaN
sum(is.nan(df1[,1]))
#Test if any INF
sum(is.finite(df1$id))
#Test if any NA's
sum(is.na(df2))
#Find the col/row with the NA
res = is.na(df2)
> colSums(res)
      id      age weight
      0       0       1
> rowSums(res)
[1] 0 0 0 1 0 0
```



Logical operators

- ▶ Useful for chaining operations

- ▶ | (OR), & (AND), ! (NOT)

- ▶ for instance, you want values greater than 10 or less than 5

```
df[ df[,1] > 10 | df[,1] < 5]
```

```
df[ !is.na(df[,1])]
```

```
df[ df[,2] == "gene" & df[,1] > 10]
```

- ▶ == is the equal operator (not =, which assigns)



Flow control

► If statements for branching

```
if (x == 3) {  
    doSomething()  
} else {  
    doSomethingElse()  
}
```

► loops

► slow in R

```
a = seq(1,10)  
for(i in 1:length(a)){  
    cat(a[i], "\n")  
}
```



Apply

- ▶ R operations are usually vectorized, so loops are avoided.
- ▶ Loops are quite slow in R

```
mat = matrix(seq(1,10), 2)
# how to get rowwise or columnwise mean?
# could use a loop, but you have to specify dimension
# margin 1 = row, 2 = cols
apply(mat, 1, mean)
apply(mat, 2, mean)
```



Sorting

- ▶ **Sorting** returns an array in sorted order

```
> b = c(1, 9, 8, 2, 4, 6)
```

```
> sort(b, decreasing=F)
```

- ▶ **Ordering** returns a vector of array indices that puts the vector in a particular order

- ▶ more useful for sorting data frames

```
> b = c(1, 9, 8, 2, 4, 6)
```

```
b.order = order(b, decreasing=F)
```

```
> b.order
```

```
[1] 1 4 5 6 3 2
```

```
> b[b.order]
```

```
[1] 1 2 4 6 8 9
```



Reading in a script

- ▶ Suppose you've written a R script that processes your data
- ▶ `source()` will read it in as if you typed it and execute all commands listed in the file
- ▶ `source("R_script_file.R")`



Additional resources

- ▶ Many resources but sometimes difficult to access
 - ▶ R-bloggers.com
 - ▶ R-project.org – FAQ and manuals
 - ▶ bioconductor.org
 - ▶ UCLA ATS (<http://www.ats.ucla.edu/stat/r/>)
 - ▶ Quick R <http://www.statmethods.net/index.html>
 - ▶ <http://onertipaday.blogspot.com/>
 - ▶ <http://wiki.stdout.org/rcookbook/>
 - ▶ <http://gettinggeneticsdone.blogspot.com/>
 - ▶ manuals.bioinformatics.ucr.edu/workshops
- ▶ Help sites
 - ▶ stackoverflow.com
 - ▶ nabble.com
 - ▶ google
- ▶ Good free books
 - ▶ An Introduction to R (on r-project.org)
 - ▶ Using R



Additional resources

▶ Courses

▶ Coursera

▶ <https://class.coursera.org/rprog-002>

▶ EdX.org

▶ https://www.edx.org/course/harvardx/harvardx-ph525x-data-analysis-genomics-1401#.U0xFnvldW_g

▶ Stanford Class

▶ <https://class.stanford.edu/courses/HumanitiesScience/StatLearning/Winter2014/about>

▶ Google YouTube

▶ <https://www.youtube.com/watch?v=iffR3fVv4xw&list=PLOU2XLYxmslK9qQfztXeybpHvru-TrqAP>

▶ Bioconductor

▶ <http://www.bioconductor.org/help/events/>



Shortcuts

- ▶ To interrupt

- ▶ MAC: <CMD> <Option> <.> or <ESC>
- ▶ PC/Linux: Ctrl-C

- ▶ sessionInfo()

- ▶ this displays details about your environment, what packages are loaded, etc... useful for debugging



Workshop 3: Introduction to R

▶ Day 2

▶ Statistical methods

- ▶ Probability distributions
- ▶ Random number generation
- ▶ Common tests
 - Fisher's exact test, chi square, qvalue, etc...

▶ Visualization

- ▶ Plotting functions
- ▶ Customizing plots
- ▶ Saving plots for publication



Workshop 3: Introduction to R

▶ **Day 3**

- ▶ Automating or scripting R
 - ▶ Batch mode
- ▶ Packages that extend R
 - ▶ Installation
 - ▶ Documentation
 - ▶ Useful packages
- ▶ Bioconductor
 - ▶ Extending R for the life sciences
 - ▶ Useful bioconductor packages
 - ▶ Some basic bioC concepts
 - ▶ Bioconductor demo

